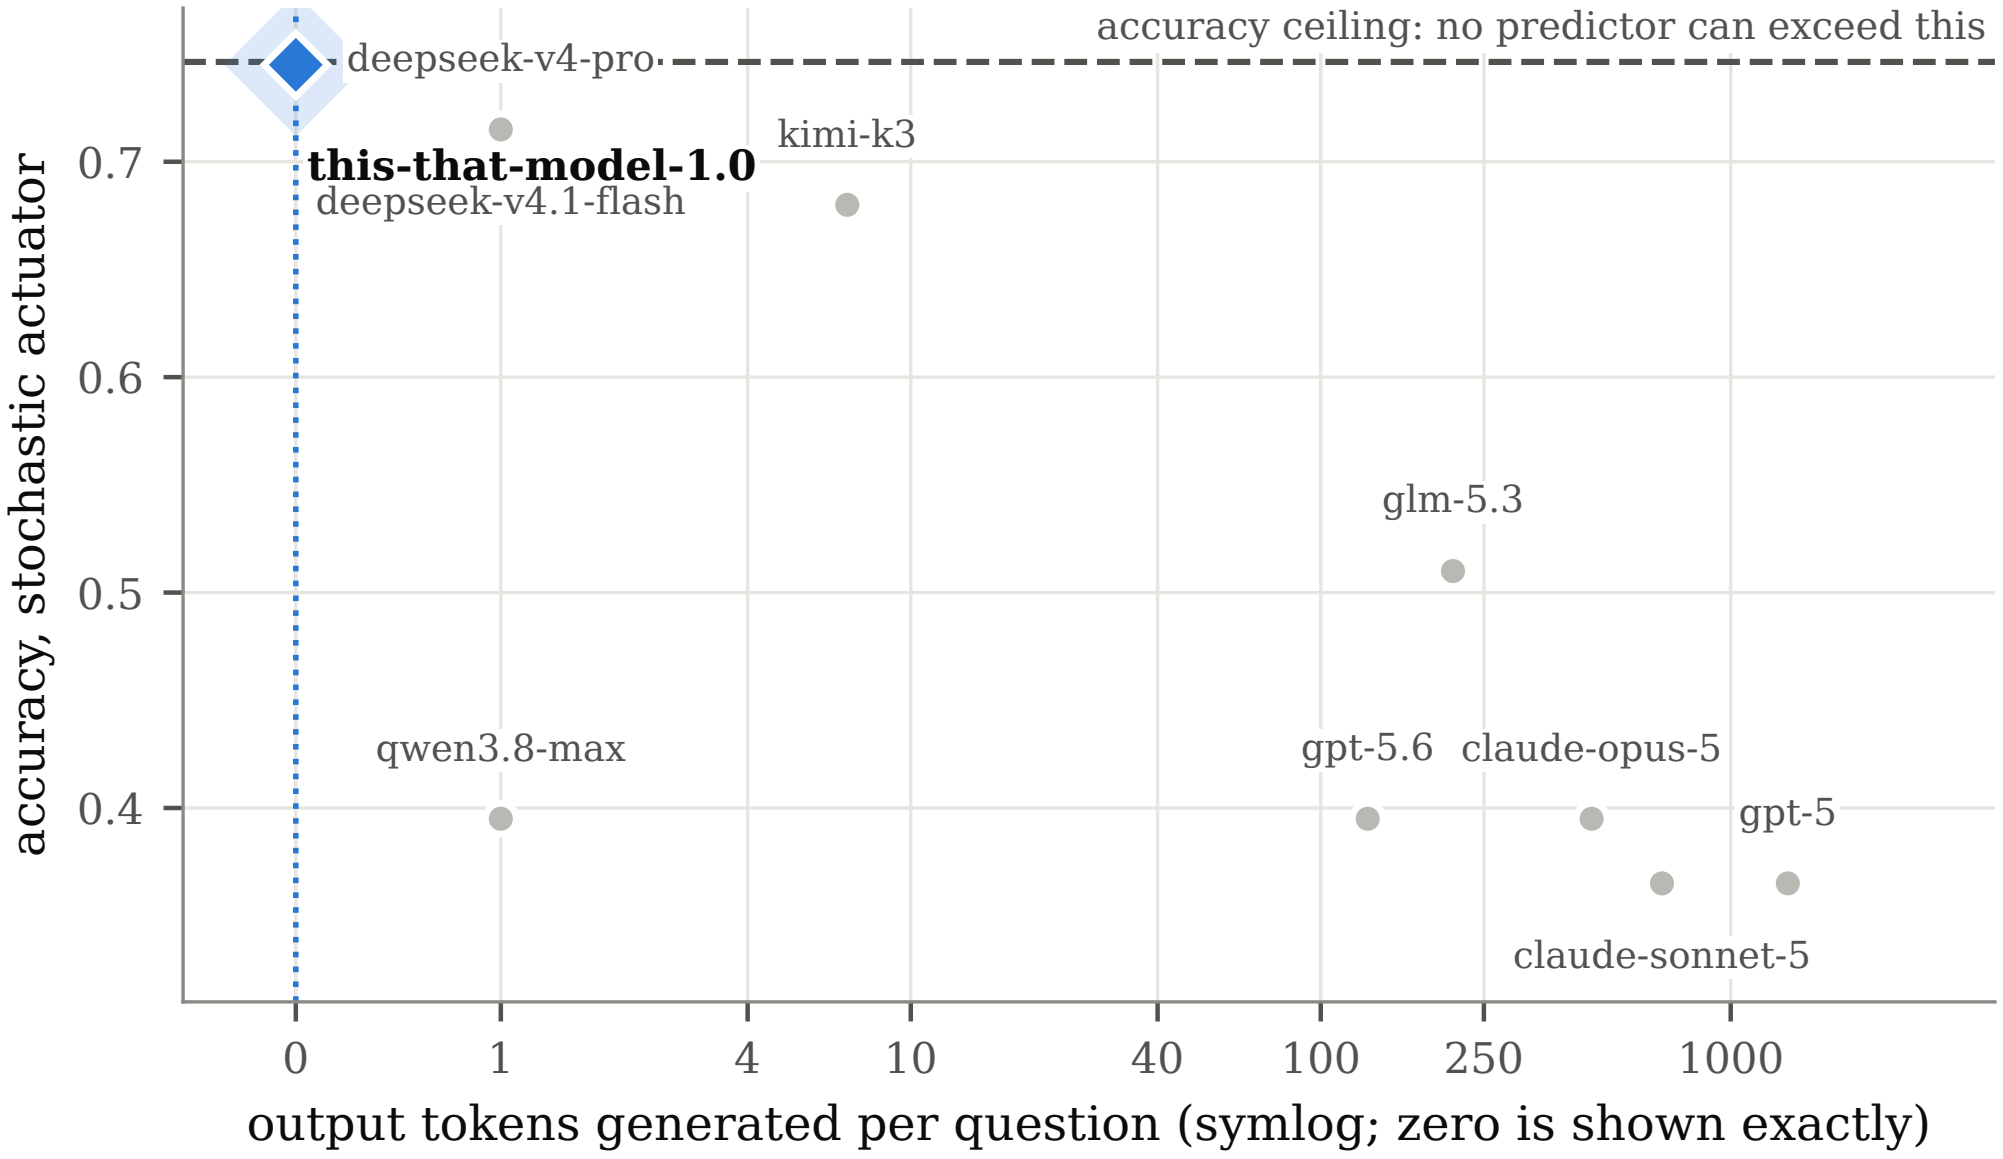
accuracy ceiling: no predictor can exceed this
deepseek-v4-pro
this-that-model-1.0
deepseek-v4.1-flash
kimi-k3
glm-5.3
qwen3.8-max
gpt-5.6
claude-opus-5
gpt-5
claude-sonnet-5
accuracy, stochastic actuator
0.7
0.6
0.5
0.4
0
1
4
10
40
100
250
1000
output tokens generated per question (symlog; zero is shown exactly)