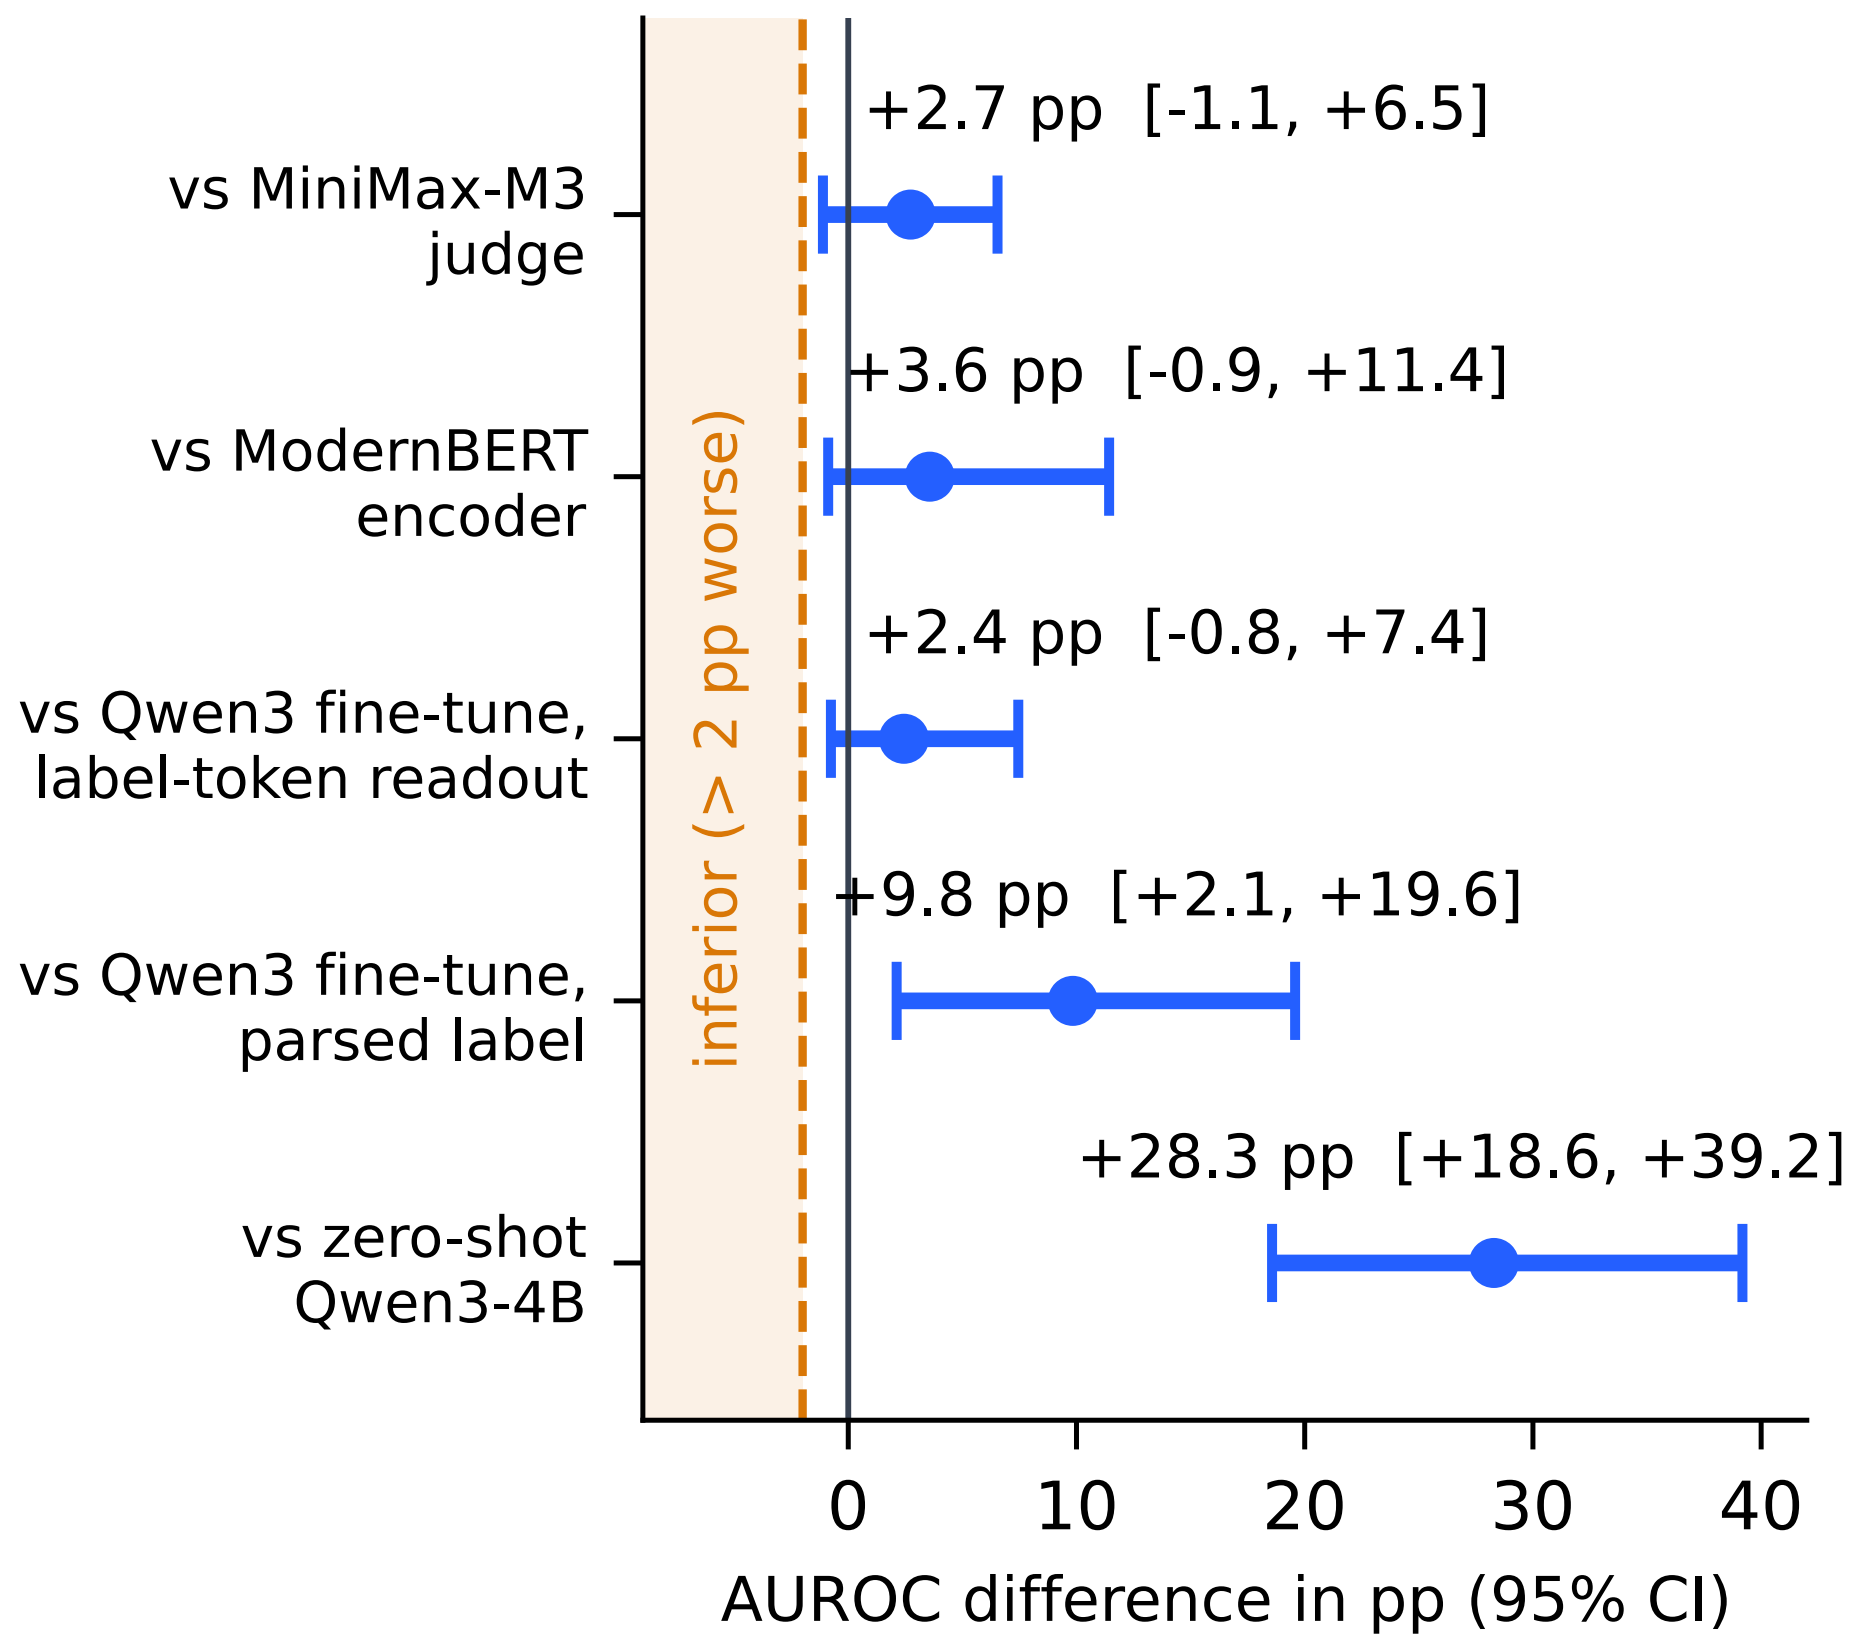
+2.7 pp [-1.1, +6.5]
vs MiniMax-M3 judge
+3.6 pp [-0.9, +11.4]
vs ModernBERT encoder
+2.4 pp [-0.8, +7.4]
vs Qwen3 fine-tune, label-token readout
+9.8 pp [+2.1, +19.6]
vs Qwen3 fine-tune, parsed label
+28.3 pp [+18.6, +39.2]
vs zero-shot Qwen3-4B
inferior (> 2 pp worse)
0
10
20
30
40
AUROC difference in pp (95% CI)