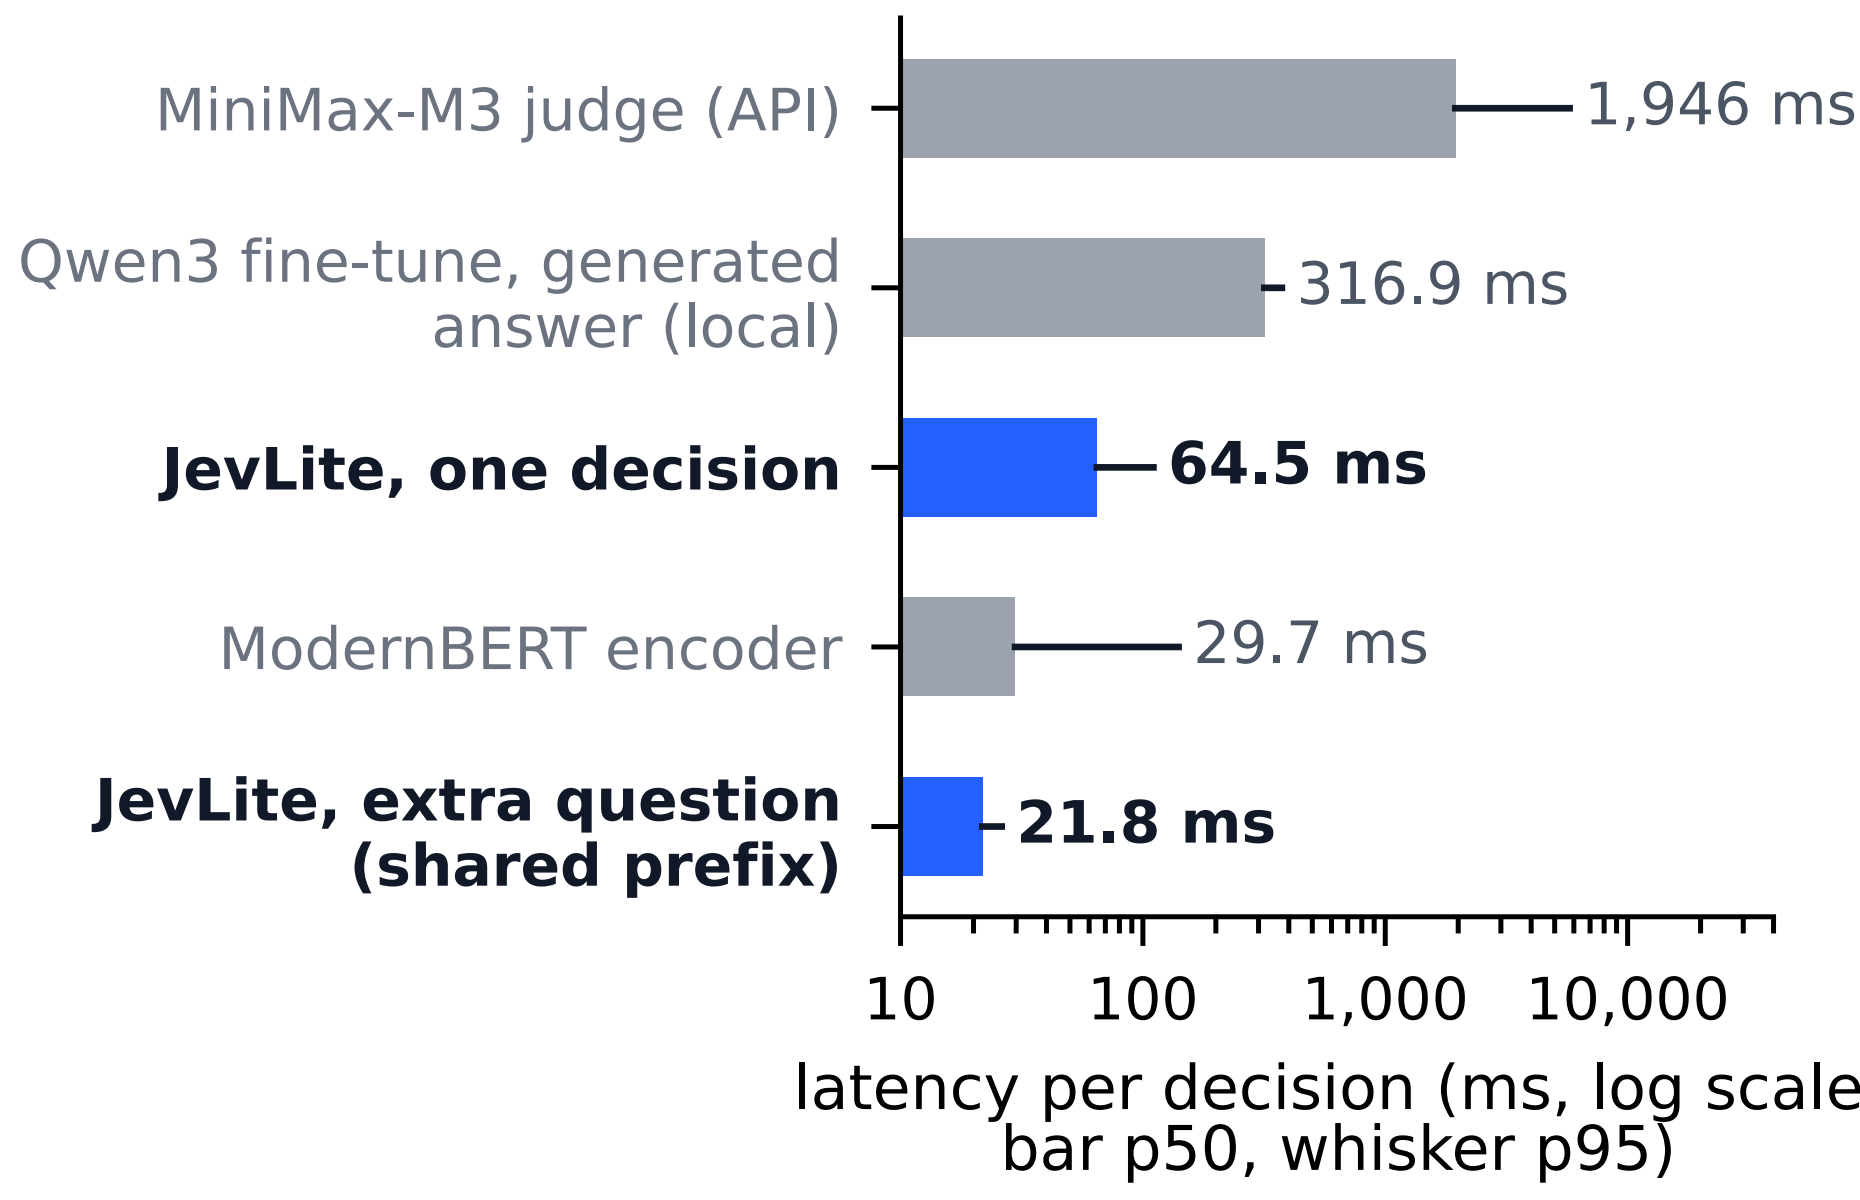
MiniMax-M3 judge (API)
1,946 ms
Qwen3 fine-tune, generated answer (local)
316.9 ms
JevLite, one decision
64.5 ms
ModernBERT encoder
29.7 ms
JevLite, extra question (shared prefix)
21.8 ms
10
100
1,000
10,000
latency per decision (ms, log scale
bar p50, whisker p95)