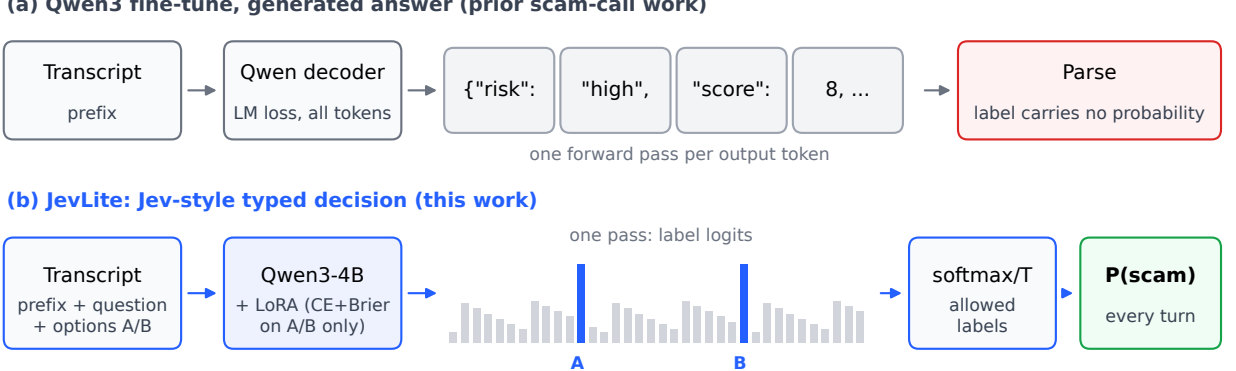
(a) Qwen3 fine-tune, generated answer (prior scam-call work)
Transcript
prefix
Qwen decoder
LM loss, all tokens
{"risk":
"high",
"score":
8, ...
Parse
label carries no probability
one forward pass per output token
(b) JevLite: Jev-style typed decision (this work)
Transcript
prefix + question
+ options A/B
Qwen3-4B
+ LoRA (CE+Brier
on A/B only)
one pass: label logits
A
B
softmax/T
allowed
labels
P(scam)
every turn