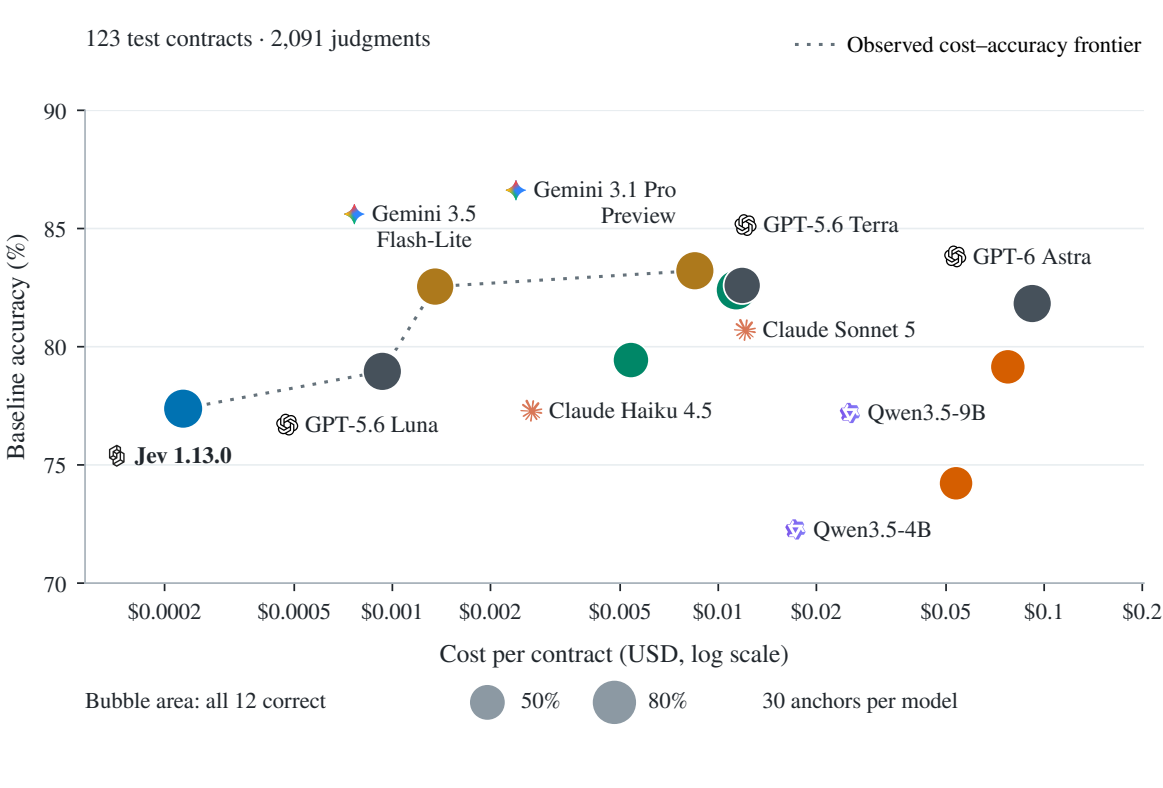
123 test contracts · 2,091 judgments
Observed cost–accuracy frontier
Baseline accuracy (%)
90
85
80
75
70
Gemini 3.1 Pro Preview
Gemini 3.5 Flash-Lite
GPT-5.6 Terra
GPT-6 Astra
Claude Sonnet 5
Claude Haiku 4.5
Qwen3.5-9B
GPT-5.6 Luna
Jev 1.13.0
Qwen3.5-4B
$0.0002
$0.0005
$0.001
$0.002
$0.005
$0.01
$0.02
$0.05
$0.1
$0.2
Cost per contract (USD, log scale)
Bubble area: all 12 correct
50%
80%
30 anchors per model